

Cost models and advanced Futhark programming

Troels Henriksen (athas@sigkill.dk)
Some material by Martin Elsman

DIKU
University of Copenhagen

20th of November, 2019

Agenda

Parallel cost models

Prefix sums (scans)

Using scans

Auxiliary

Parallel cost models

Prefix sums (scans)

Using scans

Auxiliary

The need for cost models

Which is better?

```
import numpy as np

def inc_scalar(x):
    for i in range(len(x)):
        x[i] = x[i] + 1

def inc_par(x):
    return x + np.ones(x.shape)
```

The need for cost models

Which is better?

```
import numpy as np

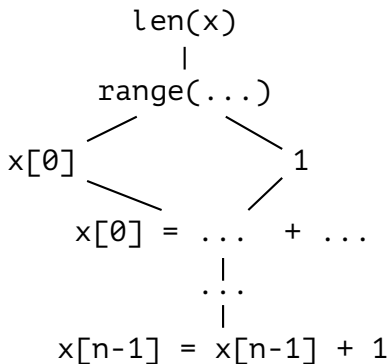
def inc_scalar(x):
    for i in range(len(x)):
        x[i] = x[i] + 1

def inc_par(x):
    return x + np.ones(x.shape)
```

Intuitively, `inc_par` is better because it is “more parallel”.


Parallel cost models make this notion precise.

Dependency DAG for `inc_scalar`



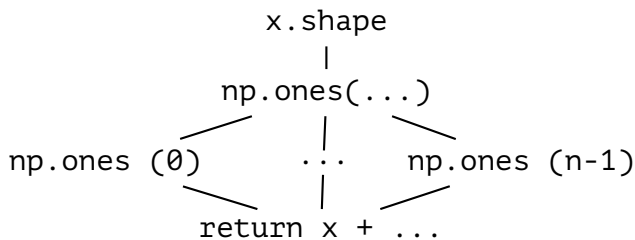
- Total count of nodes is the *work*, $W(p)$.
- Length of longest path from a leaf to the root is the *span*.
- **With an infinite number of processors, if a program p has span k , written $S(p) = k$, the program can execute in $O(k)$ time.**
- Here, $W(p) = O(n)$, $S(p) = O(n)$.

Dependency DAG for `inc_par`



What is the work and span complexity?

Dependency DAG for `inc_par`



What is the work and span complexity?

- $W(p) = O(n)$
- $S(p) = O(1)$

Parallel cost model based on work and span

Instead of giving just a simple cost-model based on the total notion of work carried out by a program, we give instead a *refined* cost model, which aims at providing both:

- a notion of how much total work (W) the program does;
- a notion of the *span*¹ (S) of the program, specifying the maximum depth required by the computation.

Notice:

- The span is the length of the longest sequence of operations that must be performed sequentially due to data dependencies.
- With an infinite number of processors, if a program p has span k , written $S(p) = k$, the program can execute in $O(k)$ time.

¹Sometimes also called *depth*.

Writing T_i for the time taken to execute an algorithm on i processors, Brent's Theorem states that

$$\frac{T_1}{p} \leq T_p \leq T_\infty + \frac{T_1}{p}$$

Proof sketch: At level j of the DAG there are M_j independent operations, which can clearly be executed by p processors in time

$$\left\lceil \frac{M_j}{p} \right\rceil$$

Sum these for each level of the DAG.



Ramification: We can simulate an “infinitely parallel” machine on a real machine at an overhead proportional to the amount of “missing” hardware parallelism.

Language-based cost models

- Tallying up levels in an infinite DAG is impractical for real programs. Instead we prefer a *language-based cost model*
- E.g. $W(x + y)$ is defined as $W(x) + W(y)$.
- The following slides define work and span cost for a small subset of Futhark.
- Write $\llbracket e \rrbracket$ for the result of evaluating expression e (we are being intuitive about scopes and such).

Language-based cost models

- Tallying up levels in an infinite DAG is impractical for real programs. Instead we prefer a *language-based cost model*
- E.g. $W(x + y)$ is defined as $W(x) + W(y)$.
- The following slides define work and span cost for a small subset of Futhark.
- Write $\llbracket e \rrbracket$ for the result of evaluating expression e (we are being intuitive about scopes and such).

Cost model must be implementable

A provable time and space efficient implementation of NESL—Guy Blelloch and John Greiner, 1996

Simple cases

$$W(v) =$$

$$S(v) =$$

$$W(e_1 \oplus e_2) =$$

$$S(e_1 \oplus e_2) =$$

$$W(\backslash x \rightarrow e) =$$

$$S(\backslash x \rightarrow e) =$$

Simple cases

$$W(v) = 1$$

$$S(v) = 1$$

$$W(e_1 \oplus e_2) =$$

$$S(e_1 \oplus e_2) =$$

$$W(\backslash x \rightarrow e) =$$

$$S(\backslash x \rightarrow e) =$$

Simple cases

$$W(v) = 1$$

$$S(v) = 1$$

$$W(e_1 \oplus e_2) = W(e_1) + W(e_2) + 1$$

$$S(e_1 \oplus e_2) = S(e_1) + S(e_2) + 1$$

$$W(\backslash x \rightarrow e) =$$

$$S(\backslash x \rightarrow e) =$$

Simple cases

$$W(v) = 1$$

$$S(v) = 1$$

$$W(e_1 \oplus e_2) = W(e_1) + W(e_2) + 1$$

$$S(e_1 \oplus e_2) = S(e_1) + S(e_2) + 1$$

$$W(\backslash x \rightarrow e) = 1$$

$$S(\backslash x \rightarrow e) = 1$$

Simple cases

$$W(v) = 1$$

$$S(v) = 1$$

$$W(e_1 \oplus e_2) = W(e_1) + W(e_2) + 1$$

$$S(e_1 \oplus e_2) = S(e_1) + S(e_2) + 1$$

$$W(\backslash x \rightarrow e) = 1$$

$$S(\backslash x \rightarrow e) = 1$$

$$W([e_1, \dots, e_n]) =$$

$$S([e_1, \dots, e_n]) =$$

$$W((e_1, \dots, e_n)) =$$

$$S((e_1, \dots, e_n)) =$$

Implications?

Simple cases

$$W(v) = 1$$

$$S(v) = 1$$

$$W(e_1 \oplus e_2) = W(e_1) + W(e_2) + 1$$

$$S(e_1 \oplus e_2) = S(e_1) + S(e_2) + 1$$

$$W(\backslash x \rightarrow e) = 1$$

$$S(\backslash x \rightarrow e) = 1$$

$$W([e_1, \dots, e_n]) = W(e_1) + \dots + W(e_n) + 1$$

$$S([e_1, \dots, e_n]) = S(e_1) + \dots + S(e_n) + 1$$

$$W((e_1, \dots, e_n)) = W(e_1) + \dots + W(e_n) + 1$$

$$S((e_1, \dots, e_n)) = S(e_1) + \dots + S(e_n) + 1$$

Implications?

Interesting cases

$W(\text{iota } e) =$

$S(\text{iota } e) =$

Interesting cases

$$W(\text{iota } e) = W(e) + \llbracket e \rrbracket$$

$$S(\text{iota } e) = S(e) + 1$$

Interesting cases

$$W(\text{iota } e) = W(e) + \llbracket e \rrbracket$$

$$S(\text{iota } e) = S(e) + 1$$

$$W(\text{let } x = e \text{ in } e') =$$

$$S(\text{let } x = e \text{ in } e') =$$

Interesting cases

$$W(\text{iota } e) = W(e) + \llbracket e \rrbracket$$

$$S(\text{iota } e) = S(e) + 1$$

$$W(\text{let } x = e \text{ in } e') = W(e) + W(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

$$S(\text{let } x = e \text{ in } e') = S(e) + S(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

Interesting cases

$$W(\text{iota } e) = W(e) + \llbracket e \rrbracket$$

$$S(\text{iota } e) = S(e) + 1$$

$$W(\text{let } x = e \text{ in } e') = W(e) + W(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

$$S(\text{let } x = e \text{ in } e') = S(e) + S(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

$$W(e_1 \ e_2) =$$

$$S(e_1 \ e_2) =$$

Interesting cases

$$W(\text{iota } e) = W(e) + \llbracket e \rrbracket$$

$$S(\text{iota } e) = S(e) + 1$$

$$W(\text{let } x = e \text{ in } e') = W(e) + W(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

$$S(\text{let } x = e \text{ in } e') = S(e) + S(e'[x \mapsto \llbracket e \rrbracket]) + 1$$

$$W(e_1 \ e_2) = W(e_1) + W(e'[x \mapsto \llbracket e_2 \rrbracket]) + 1$$

$$\text{where } \llbracket e_1 \rrbracket = \backslash x \rightarrow e'$$

$$S(e_1 \ e_2) = S(e_1) + S(e'[x \mapsto \llbracket e_2 \rrbracket]) + 1$$

$$\text{where } \llbracket e_1 \rrbracket = \backslash x \rightarrow e'$$

Work and span for map

$$W(\text{map } e_1 \ e_2) =$$

$$S(\text{map } e_1 \ e_2) =$$

Work and span for map

$$\begin{aligned} W(\text{map } e_1 \ e_2) &= \\ &W(e_1) + W(e_2) + W(e'[x \mapsto v_1]) + \dots + W(e'[x \mapsto v_n]) \\ \text{where } \llbracket e_1 \rrbracket &= \backslash x \rightarrow e' \\ \text{where } \llbracket e_2 \rrbracket &= [v_1, \dots, v_n] \end{aligned}$$

$$\begin{aligned} S(\text{map } e_1 \ e_2) &= \\ &S(e_1) + S(e_2) + \max(S(e'[x \mapsto v_1]), \dots, S(e'[x \mapsto v_n])) + 1 \\ \text{where } \llbracket e_1 \rrbracket &= \backslash x \rightarrow e' \\ \text{where } \llbracket e_2 \rrbracket &= [v_1, \dots, v_n] \end{aligned}$$

Reduction by contraction

```
let npow2 (n:i32) : i32 =  
  loop a = 2 while a < n do 2*a  
  
-- Pad a vector to make its size a power of two  
let padpow2 [n] (ne: i32) (v:[n]i32) : []i32 =  
  concat v (replicate (npow2 n - n) ne)  
  
-- Reduce by contraction  
let red (xs : []i32) : i32 =  
  let xs =  
    loop xs = padpow2 0 xs  
    while length xs > 1 do  
      let n = length xs / 2  
      in map2 (+) xs[0:n] xs[n:2*n]  
  in xs[0]
```

Work and span of `loop`

$$W(\text{loop } x = e_1 \text{ while } e_2 \text{ do } e_3) =$$

$$S(\text{loop } x = e_1 \text{ while } e_2 \text{ do } e_3) =$$

Work and span of `loop`

$W(\text{loop } x = e_1 \text{ while } e_2 \text{ do } e_3) =$
if $\llbracket e_2 \rrbracket = \text{true}$
then $1 + W(e_2)$
else $1 + W(e_2) + W(\text{loop } x = \llbracket e_3[x \mapsto e_1] \rrbracket \text{ while } e_2 \text{ do } e_3)$

$S(\text{loop } x = e_1 \text{ while } e_2 \text{ do } e_3) =$
if $\llbracket e_2 \rrbracket = \text{true}$
then $1 + S(e_2)$
else $1 + S(e_2) + S(\text{loop } x = \llbracket e_3[x \mapsto e_1] \rrbracket \text{ while } e_2 \text{ do } e_3)$

Work and Span for $n^{\text{pow}2}$

Work and Span for $\text{npow2 } n$

By inspection, we have

$$W(\text{npow2 } n) = S(\text{npow2 } n) = O(\log n)$$

Work and Span for $\text{padpow2 } ne \text{ } v$

Work and Span for npow2 n

By inspection, we have

$$W(\text{npow2 } n) = S(\text{npow2 } n) = O(\log n)$$

Work and Span for padpow2 ne v

Because $\text{npow2 } n \leq 2n$, we have (where $n = \text{length } v$)

$$\begin{aligned} W(\text{padpow2 } ne \ v) \\ &= W(\text{concat } v \ (\text{replicate } (\text{npow2 } n - n) \ ne)) \\ &= O(n) \end{aligned}$$

$$S(\text{padpow2 } ne \ v) = O(\log n)$$

Work and Span for red

Work and Span for npow2 n

By inspection, we have

$$W(\text{npow2 } n) = S(\text{npow2 } n) = O(\log n)$$

Work and Span for padpow2 ne v

Because $\text{npow2 } n \leq 2n$, we have (where $n = \text{length } v$)

$$\begin{aligned}
W(\text{padpow2 } ne \ v) &= W(\text{concat } v \ (\text{replicate } (\text{npow2 } n - n) \ ne)) \\
&= O(n)
\end{aligned}$$

$$S(\text{padpow2 } ne \ v) = O(\log n)$$

Work and Span for red

Each loop iteration in red has span $O(1)$. Because the loop is iterated at-most $\log(2n)$ times, we have (where $n = \text{length } v$)

$$\begin{aligned}
W(\text{red } v) &= O(n) + O(n/2) + O(n/4) + \cdots + O(1) = O(n) \\
S(\text{red } v) &= O(\log n)
\end{aligned}$$

Work efficiency

A parallel algorithm is said to be *work efficient* if it has at most the same work as the best sequential algorithm.

Is red work efficient?

Work efficiency

A parallel algorithm is said to be *work efficient* if it has at most the same work as the best sequential algorithm.

Is red work efficient?

Yes, because it does $O(n)$ work, which is as good as a sequential summation.

Is it also *efficient*?

Performance Compared to the Built-in Reduction SOAC

```
-- ==  
-- entry: test_red test_reduce  
-- random input { [10000000]i32 }  
entry test_red = red  
entry test_reduce = reduce (+) 0
```

Performance Compared to the Built-in Reduction SOAC

```
-- ==  
-- entry: test_red test_reduce  
-- random input { [10000000]i32 }  
entry test_red = red  
entry test_reduce = reduce (+) 0  
  
$ futhark bench --backend=opencl reduce.fut  
Compiling reduce.fut...  
Results for reduce.fut:test_red:  
dataset [10000000]i32:      4675.40 $\mu$ s  
Results for reduce.fut:test_reduce:  
dataset [10000000]i32:      273.80 $\mu$ s
```

Performance Compared to the Built-in Reduction SOAC

```
-- ==  
-- entry: test_red test_reduce  
-- random input { [10000000]i32 }  
entry test_red = red  
entry test_reduce = reduce (+) 0
```

```
$ futhark bench --backend=opencl reduce.fut  
Compiling reduce.fut...  
Results for reduce.fut:test_red:  
dataset [10000000]i32:      4675.40 $\mu$ s  
Results for reduce.fut:test_reduce:  
dataset [10000000]i32:      273.80 $\mu$ s
```

If you are not using `futhark bench`, then you are probably doing it wrong.

Parallel cost models

Prefix sums (scans)

Using scans

Auxiliary

Inclusive and exclusive prefix sum

Exclusive prefix sum ("prescan")

Given

[1, 2, 3, 4]

produce

[0, 1, 3, 6]

Inclusive prefix sum

Given

[1, 2, 3, 4]

produce

[1, 3, 6, 10]

Prefix sums are scans

Generalising the addition and zero used by a prefix sum to an arbitrary associative operator \oplus and neutral element 0_{\oplus} , we get *scan*.

-- The scan in Futhark is inclusive.

```
> scan (+) 0 [1,2,3,4]  
[1, 3, 6, 10]
```

Prefix sums are scans

Generalising the addition and zero used by a prefix sum to an arbitrary associative operator \oplus and neutral element 0_{\oplus} , we get *scan*.

-- The scan in Futhark is inclusive.

```
> scan (+) 0 [1,2,3,4]  
[1, 3, 6, 10]
```

- Scans are a fundamental building block on parallelising seemingly-sequential algorithms.
- Let us see how scans can be computed in parallel.

Sequential prefix sum

```
acc = 0
for i < n:
    acc = acc + input[i]
    scanned[i] = acc
```

Sequential prefix sum

```
acc = 0
for i < n:
    acc = acc + input[i]
    scanned[i] = acc
```

Work: $O(n)$

Span: $O(n)$

Brute force

To calculate the prefix sum of $[x_0, \dots, x_{n-1}]$, compute

$$\begin{aligned} &[sum([x_0]) \\ &sum([x_0, x_1]) \\ &\vdots \\ &sum([x_0, x_1, \dots, x_{n-1}])] \end{aligned}$$

Assume $S(sum([x_0, \dots, x_{n-1}])) = \log_2(n)$.

Brute force

To calculate the prefix sum of $[x_0, \dots, x_{n-1}]$, compute

$$\begin{aligned} & \text{sum}([x_0]) \\ & \text{sum}([x_0, x_1]) \\ & \vdots \\ & \text{sum}([x_0, x_1, \dots, x_{n-1}]) \end{aligned}$$

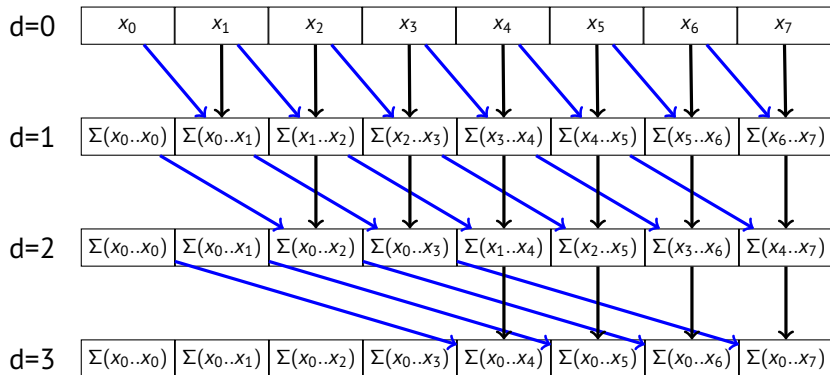
Assume $S(\text{sum}([x_0, \dots, x_{n-1}])) = \log_2(n)$.

Work: $O(\sum_{i < n} i) = O(n^2)$

Span: $O(\max(S(\text{sum}([x_0])), \dots, S(\text{sum}([x_0, \dots, x_{n-1}])))) = O(\log_2(n))$

Terrible. The sequential implementation is faster for large n !

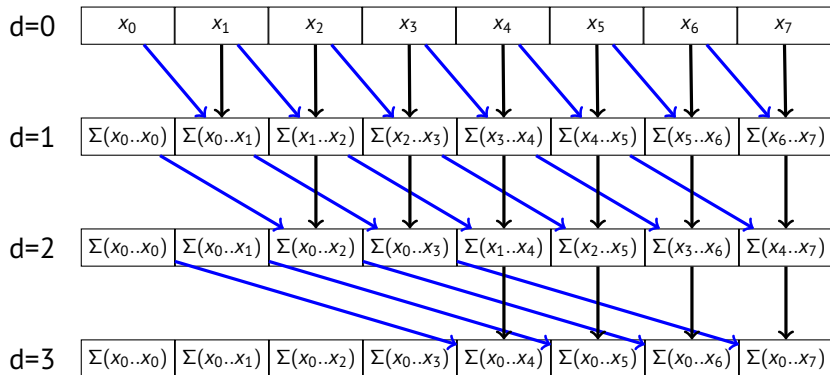
Hillis–Steele scan (1986)²



For each d , element x_i^d is updated by $x_{i-2^d}^{d-1} + x_i^{d-1}$.

²This slide took 60 minutes to make.

Hillis–Steele scan (1986)²



For each d , element x_i^d is updated by $x_{i-2^d}^{d-1} + x_i^{d-1}$.

Work: For $n = 2^m$, $O(\sum_{i < m} 2^m - 2^i) = O(n \log(n))$

Span: $\log(n)$

²This slide took 60 minutes to make.

Work-efficient scan

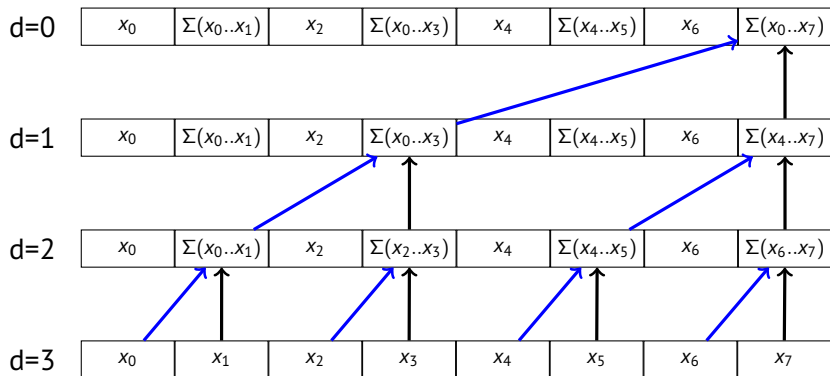
Two passes

Up-sweep Build a balanced binary tree of partial sums stored in every other cell.

Down-sweep Use the partial sums to fill out the missing parts.

The binary tree does not actually exist as a recursive pointer structure, but is just a communications concept.


Up-sweep (“reduction phase”)³



$$x_i^d = x_{i-2^{m-d-1}}^{d+1} + x_i^{d+1}$$

³This slide took 30 minutes to make.

Up-sweep (“reduction phase”)³



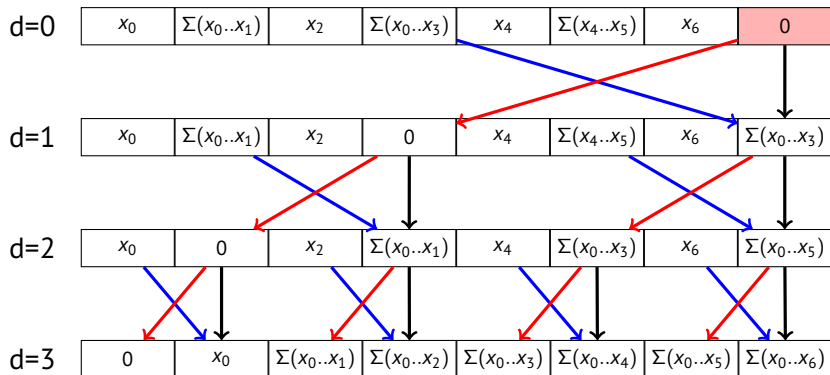
$$x_i^d = x_{i-2^{m-d-1}}^{d+1} + x_i^{d+1}$$

Work: For $n = 2^m$, $O(\sum_{i < m} 2^i) = O(n)$

Span: $\log(n)$

³This slide took 30 minutes to make.

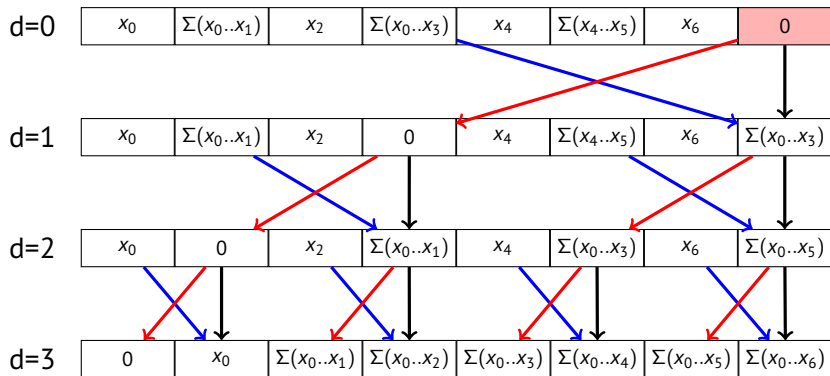
Down-sweep⁴



Inverse indexing of the up-sweep phase.

⁴This slide took 15 minutes to make.

Down-sweep⁴



Inverse indexing of the up-sweep phase.

Work: For $n = 2^m$, $O(\sum_{i \leq m} 2^i) = O(n)$

Span: $\log(n)$

⁴This slide took 15 minutes to make.

Work efficient scan

Complexity of *scan* on size- n input

Work: $O(n)$

Span: $\log(n)$

- Optimal, as *reduce* is the same.
- Can now depend on scan as a relatively cheap building block.

Real-world scan implementations are often very different for technical reasons, but we can depend on these asymptotics when analysing and designing parallel algorithms.

Parallel cost models

Prefix sums (scans)

Using scans

Auxiliary

Filtering

Suppose we wish to remove negative elements from the list

let as = [-1, 2, -3, 4, 5, -6]

Filtering

Suppose we wish to remove negative elements from the list

```
let as = [-1, 2, -3, 4, 5, -6]
```

For each element, see if we want to keep it:

```
let keep = map (\a -> if a >= 0 then 1 else 0) as  
-- [ 0, 1, 0, 1, 1, 0]
```

Filtering

Suppose we wish to remove negative elements from the list

```
let as = [-1, 2, -3, 4, 5, -6]
```

For each element, see if we want to keep it:

```
let keep = map (\a -> if a >= 0 then 1 else 0) as  
-- [ 0, 1, 0, 1, 1, 0]
```

```
let offsets1 = scan (+) 0 keep  
-- [ 0, 1, 1, 2, 3, 3]
```

Filtering

Suppose we wish to remove negative elements from the list

```
let as = [-1, 2, -3, 4, 5, -6]
```

For each element, see if we want to keep it:

```
let keep = map (\a -> if a >= 0 then 1 else 0) as  
-- [ 0, 1, 0, 1, 1, 0]
```

```
let offsets1 = scan (+) 0 keep  
-- [ 0, 1, 1, 2, 3, 3]
```

```
let offsets = map (\x -> x - 1) offsets1  
-- [-1, 0, 0, 1, 2, 2]
```

Filtering

Suppose we wish to remove negative elements from the list

```
let as = [-1, 2, -3, 4, 5, -6]
```

For each element, see if we want to keep it:

```
let keep = map (\a -> if a >= 0 then 1 else 0) as  
-- [ 0, 1, 0, 1, 1, 0]
```

```
let offsets1 = scan (+) 0 keep  
-- [ 0, 1, 1, 2, 3, 3]
```

```
let offsets = map (\x -> x - 1) offsets1  
-- [-1, 0, 0, 1, 2, 2]
```

`offsets[i]` now indicates position in filtered list iff `keep[i] == 1`.

scatter

scatter xs is vs computes equivalent of the imperative pseudocode

```
for j < n:  
    xs[is[j]] = vs[j]
```

- Out-of-bound writes are ignored
- Writing different values to same index is *undefined*⁵
- Work $O(n)$, span $O(1)$

Just what we need for filtering!

⁵reduce_by_index handles conflicts with provided operator.

scatter

scatter xs is vs computes equivalent of the imperative pseudocode

```
for j < n:  
  xs[is[j]] = vs[j]
```

- Out-of-bound writes are ignored
- Writing different values to same index is *undefined*⁵
- Work $O(n)$, span $O(1)$

Just what we need for filtering!

```
scatter (replicate (last offsets1) 0)  
  (map2 (\i k -> if k == 1 then i else -1)  
    offsets keep)  
  as
```

⁵reduce_by_index handles conflicts with provided operator.

Implementing filter

```
let filter 'a (p: a -> bool) (as: []a): []a =  
  let keep = map (\a -> if p a then 1 else 0) as  
  let offsets1 = scan (+) 0 keep  
  let num_to_keep = reduce (+) 0 keep  
  in if num_to_keep == 0  
    then []  
    else scatter (replicate num_to_keep as[0])  
      (map2 (\i k ->  
        if k == 1  
        then i-1  
        else -1)  
        offsets1 keep)  
    as
```


Radix sort

- Many classical sorting algorithms are a poor fit for data parallelism, but *radix sort* works well.
- Radix-2 sort works by repeatedly partitioning elements according to one bit at a time, while preserving the ordering of the previous steps.

Example with radix-10

3 2 6		6 9 0		7 0 4		3 2 4
4 5 3		7 5 1		6 0 8		4 3 8
6 0 8		4 5 3		3 2 6		4 5 6
8 3 5		7 0 4		8 3 5		6 0 5
7 5 1	⇒	8 3 5		4 3 5		6 9 5
4 3 5		4 3 5	⇒	7 5 1		7 0 1
7 0 4		3 2 6		4 5 3		7 5 3
6 9 0		6 0 8		6 9 0		8 3 0

Example with radix-10



- **Radix sort is not as general as a comparison-based sort.**
- Assumes sorting key can be decomposed into “digits”.

Sorting `xs:[n]u32` by bit `b`

```
-- 1 if bit b set.  
let check_bit b x =  
    (i32.u32 (x >> u32.i32 b)) & 1
```

Sorting xs : [n]u32 by bit b

```
-- 1 if bit b set.  
let check_bit b x =  
    (i32.u32 (x >> u32.i32 b)) & 1  
  
let bits = map (check_bit b) xs  
let bits_neg = map (1-) bits  
let offs = reduce (+) 0 bits_neg
```

Sorting xs : [n]u32 by bit b

```
-- 1 if bit b set.  
let check_bit b x =  
    (i32.u32 (x >> u32.i32 b)) & 1  
  
let bits = map (check_bit b) xs  
let bits_neg = map (1-) bits  
let offs = reduce (+) 0 bits_neg
```

Example

```
b          = 0  
xs         = [0, 1, 2, 3, 4]  
bits       = [0, 1, 0, 1, 0]  
bits_neg   = [1, 0, 1, 0, 1]  
offs       = 3
```

```
let idxs0 = map2 (*)  
              bits_neg  
              (scan (+) 0 bits_neg)  
let idxs1 = map2 (*)  
              bits  
              (map (+offs) (scan (+) 0 bits))
```

```
let idxs0 = map2 (*)  
              bits_neg  
              (scan (+) 0 bits_neg)  
let idxs1 = map2 (*)  
              bits  
              (map (+offs) (scan (+) 0 bits))
```

Example

```
bits      = [0, 1, 0, 1, 0]  
bits_neg  = [1, 0, 1, 0, 1]  
offs      = 3  
idxs0     = [1, 0, 2, 0, 3]  
idxs1     = [0, 4, 0, 5, 0]  
map2 (+) idxs0 idxs1  
          = [1, 4, 2, 5, 3]
```

Then **scatter** as when **filtering**.

The whole step

```
let check_bit b x =  
  (i32.u32 (x >> u32.i32 b)) & 1  
  
let radix_sort_step [n]  
  (xs: [n]u32) (b: i32): [n]u32 =  
  let bits = map (check_bit b) xs  
  let bits_neg = map (1-) bits  
  let offs = reduce (+) 0 bits_neg  
  let idxs0 = map2 (*) bits_neg  
    (scan (+) 0 bits_neg)  
  let idxs1 = map2 (*) bits  
    (map (+offs) (scan (+) 0 bits))  
  let idxs2 = map2 (+) idxs0 idxs1  
  let idxs = map (\x->x-1) idxs2  
  let xs' = scatter (copy xs) idxs xs  
  in xs'
```


Radix sort in Futhark

```
let radix_sort [n] (xs: [n]u32): [n]u32 =  
    loop xs for i < 32 do radix_sort_step xs i
```

See worked example at <https://futhark-lang.org/examples/radix-sort.html>

Segmented scan

```
val segmented_scan [n] 't
  : (op: t -> t -> t) -> (ne: t)
  -> (flags: [n]bool) -> (as: [n]t)
  -> [n]t
```

true starts a segment and false continues a segment.

Example

```
segmented_scan (+) 0
  [true, false, true, false, false, true]
  [0, 1, 2, 3, 4, 5]
== scan (+) 0 [0,1] ++
   scan (+) 0 [2,3,4] ++
   scan (+) 0 [5]
== [0, 1, 2, 5, 9, 5]
```

Segmented reduction

```
val segmented_reduce [n] 't
  : (op: t -> t -> t) -> (ne: t)
  -> (flags: [n]bool) -> (as: [n]t)
  -> []t
```

Example

```
segmented_reduce (+) 0
  [true, false, true, false, false, true]
  [0, 1, 2, 3, 4, 5]
== reduce (+) 0 [0,1] ++
   reduce (+) 0 [2,3,4] ++
   reduce (+) 0 [5]
== [1, 9, 5]
```

Generalised histograms

Like `scatter`, but uses a provided reduce-like operator to handle multiple writes to same index.

Type

```
val reduce_by_index [k] [n] 'a :  
    (dest: *[k]a)  
    -> (f: a -> a -> a) -> (ne: a)  
    -> (is: [n]i32) -> (vs: [n]a) -> *[k]a
```

Semantics

```
for index in 0..k-1:  
    i = is[index]  
    v = vs[index]  
    dest[i] = f(as[i], v)
```

Futhark uses parallel
implementation with GPU
atomics.

Proving associativity and neutral elements

```
let op (x, i) (y, j) : (i32, i32) =  
  if x < y then (y, j) else (x, i)
```

```
let argmax [n] (xs: [n]i32) =  
  reduce op  
    (i32.smallest, -1)  
    (zip xs (iota n))
```

- Is op associative?
- Is (i32.smallest, -1) a neutral element?

argmax: associativity

```
(a 'op' b) 'op' c
== ((ax, ai) 'op' (bx, bi)) 'op' (cx, ci)
== let (x, i) = if ax < bx then (bx, bi)
                                else (ax, ai)
   in if x < cx then (cx, ci)
       else (x, i)
```

```
a 'op' (b 'op' c)
== (ax, ai) 'op' ((bx, bi) 'op' (cx, ci))
== let (x, i) = if bx < cx then (cx, ci)
                                else (bx, bi)
   in if ax < x then (x, i)
       else (ax, ai)
```

Enumerate all possible comparisons between ax, bx, and cx and show associativity for each.

E.g. for **!(ax < bx) && bx < cx && cx < ax**

```
let (x, i) = if ax < bx then (bx, bi)
              else (ax, bx)
in if x < cx then (cx, ci)
    else (x, i)
== if ax < cx then (cx, ci)
    else (ax, ai)
== (ax, ai)
```

```
let (x, i) = if bx < cx then (cx, ci)
              else (bx, bi)
in if ax < x then (x, i)
    else (ax, ai)
== if ax < cx then (cx, ci)
    else (ax, ai)
== (ax, ai)
```



argmax: neutral element

```
(a 'op' (i32.smallest, -1))  
== ((x, i) 'op' (i32.smallest, -1))  
== if x < i32.smallest then (i32.smallest, -1)  
    else (x, i)  
== (x, i)
```

```
((i32.smallest, -1) 'op' a)  
== ((i32.smallest, -1) 'op' (x, i))  
== if i32.smallest < x then (x, i)  
    else (i32.smallest, -1)  
== (x, i)
```


Commutativity?

Exercise for home: The `argmax` operator is not commutative. Try to come up with a counterexample, and see if you can change its definition such that it becomes commutative.

Commutativity?

Exercise for home: The `argmax` operator is not commutative. Try to come up with a counterexample, and see if you can change its definition such that it becomes commutative.

Commutative reductions

Futhark has a `reduce_comm` function that can be used for commutative operators. This runs faster than normal `reduce`. Not necessary for built-in operators.

Summary

- *Work* measures the total number of operations, *span* measures the longest chain of dependencies.
- Language-based cost models let us reason about program performance in a hardware-agnostic and composable way.
- Scans are a useful building block in advanced data parallel algorithms, but an efficient implementation is not straightforward.