

# Parallel Haskell

Troels Henriksen

Slides based on work by Ken Friis Larsen  
based on material by Simon Marlow

DIKU

University of Copenhagen

2nd of December, 2019

# Agenda

What is Parallel Haskell?

# All you need is $X$

- Where  $X$  is:
  - ▶ Actors, threads, transactional memory, futures, love...
- In Haskell, the approach is to give you lots of different  $X$ s

# All you need is $X$

- Where  $X$  is:
  - ▶ Actors, threads, transactional memory, futures, love...
- In Haskell, the approach is to give you lots of different  $X$ s

*“Embrace diversity (but control side effects)”*

*– Simon Peyton Jones*

# All you need is $X$

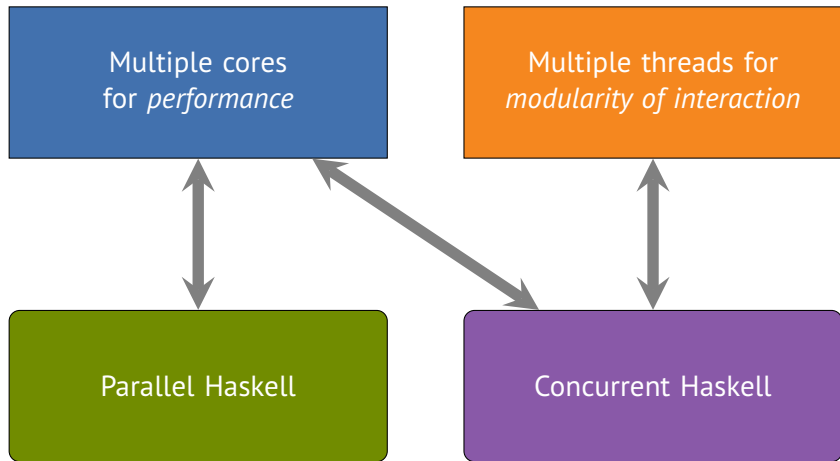
- Where  $X$  is:
  - ▶ Actors, threads, transactional memory, futures, love...
- In Haskell, the approach is to give you lots of different  $X$ s

*“Embrace diversity (but control side effects)”*

*– Simon Peyton Jones*

**What is the difference between parallelism and concurrency?**

# Parallelism vs. Concurrency



## Task parallelism

Independent tasks executed on multiple processors, possibly with very different code and data.

Primary distinguishing feature of Parallel Haskell:

- The program does “the same thing” regardless of how many cores are used (*determinism*)
- No race conditions or deadlocks
  - ▶ add parallelism without sacrificing correctness
- Parallelism is used to speed up pure (non-IO monad) Haskell code

# Lecture goals

- Understanding the GHC primitives for working with parallelism
- Tools and pitfalls
- Work with the `Par` Monad
- Use the `Par` monad for expressing dataflow parallelism



# Running example: Sudoku

- Peter Norvig's algorithm with constraint propagation. Code from the Haskell wiki
- can solve all 49,000<sup>(\*)</sup> problems in 2 mins
  - (\*) collected by Gordon Royle University of Western Australia
- input: a line of text representing a problem

				3		8	5	
	1		2					
		5		7				
	4				1			
9								
5						7	3	
	2		1					
			4				9	

4					8		5	
	3							
		7						
	2					6		
			8		4			
			1					
		6		3		7		
5		2						
1	4							

.....3.85..1.2.....5.7.....4...1...9.....5.....73..2.1.....4...9  
4.....8.5.3.....7.....2.....6.....8.4.....1.....6.3.7.5..2.....1.4.....

# Solving Sudoku problems

```
import Sudoku (solve, Grid)
import System.Environment

type Grid = ...
solve :: String -> Maybe Grid
evaluate :: a -> IO a

main :: IO ()
main = do
    [f] <- getArgs
    file <- readFile f
    let puzzles    = lines file
        solutions = map solve puzzles

    print (length (filter isJust solutions))
```

## Compile and run the program

```
$ ghc -O2 sudoku1.hs -rtsopts
```

## Compile and run the program

```
$ ghc -O2 sudoku1.hs -rtsopts  
[1 of 2] Compiling Sudoku  
[2 of 2] Compiling Main  
Linking sudoku1 ...  
$
```

```
( Sudoku.hs, Sudoku.o )  
( sudoku1.hs, sudoku1.o )
```

## Compile and run the program

```
$ ghc -O2 sudoku1.hs -rtspts  
[1 of 2] Compiling Sudoku          ( Sudoku.hs, Sudoku.o )  
[2 of 2] Compiling Main            ( sudoku1.hs, sudoku1.o )  
Linking sudoku1 ...  
$ ./sudoku1 sudoku17.1000.txt +RTS -s
```

# Compile and run the program

```
$ ghc -O2 sudoku1.hs -rtsopts
[1 of 2] Compiling Sudoku          ( Sudoku.hs, Sudoku.o )
[2 of 2] Compiling Main              ( sudoku1.hs, sudoku1.o )
Linking sudoku1 ...
$ ./sudoku1 sudoku17.1000.txt +RTS -s
    2,352,239,176 bytes allocated in the heap
    39,000,728 bytes copied during GC
    213,616 bytes maximum residency (13 sample(s))
    81,576 bytes maximum slop
    2 MB total memory in use (0 MB lost due to fragmentation)

                             Tot time (elapsed)
Gen   0           4552 colls,      0 par    0.09s   0.09s
Gen   1           13 colls,      0 par    0.00s   0.00s

INIT    time    0.00s   ( 0.00s elapsed)
MUT     time    1.99s   ( 2.01s elapsed)
GC      time    0.09s   ( 0.10s elapsed)
EXIT    time    0.00s   ( 0.00s elapsed)
Total   time    2.09s   ( 2.11s elapsed) # <- what we care about
```

## Let's use some cores

- Doing parallel computation entails specifying coordination in some way – compute  $A$  in parallel with  $B$
- This is a constraint on evaluation order
- But by design, Haskell *does not have a specified evaluation order*
- So we need to add something to the language to express constraints on evaluation order

# The Eval monad

```
import Control.Parallel.Strategies
```

```
data Eval a
```

```
instance Monad Eval
```

```
runEval :: Eval a -> a
```

```
rpar :: a -> Eval a
```

```
rseq :: a -> Eval a
```

---

<sup>1</sup>*Seq no more: Better Strategies for Parallel Haskell*



# The Eval monad

```
import Control.Parallel.Strategies
```

```
data Eval a
```

```
instance Monad Eval
```

```
runEval :: Eval a -> a
```

```
rpar :: a -> Eval a
```

```
rseq :: a -> Eval a
```

- Eval is pure
- Just for expressing sequencing between rpar/rseq
- Compositional – larger Eval sequences can be built by composing smaller ones using monad combinators
- Internal workings of Eval are simple<sup>1</sup>


---

<sup>1</sup>*Seq no more: Better Strategies for Parallel Haskell*


# What does `rpar` actually do?

- `x <- rpar e`
- `rpar` creates a *spark* by writing an entry in the *spark pool*
  - ▶ `rpar` is cheap (not a OS thread)
- The spark pool is a circular buffer
- When a processor has nothing to do, it tries to remove an entry from its own spark pool, or steal an entry from another spark pool (*work stealing*)
- When a spark is found, it is evaluated
- If the value corresponding to a spark is needed before it is computed, it *fizzles*.
- The spark pool can be full – watch out for spark overflow

# Illustration by Don Stewart



Haskell Spark Pools



Haskell Lightweight Threads



$N$  Worker (OS) Threads



CPU 1

CPU 2

CPU 3

CPU 4

## Basic Eval patterns

- To compute a in parallel with b, and return a pair of the results:

```
do a' <- rpar a  
    b' <- rseq b  
    rseq a'  
    return (a', b')
```

# Basic Eval patterns

- To compute a in parallel with b, and return a pair of the results:

```
do a' <- rpar a  
    b' <- rseq b  
    rseq a'  
    return (a', b')
```

- Alternatively:

```
do a' <- rpar a  
    b' <- rseq b  
    return (a', b')
```

# Basic Eval patterns

- To compute a in parallel with b, and return a pair of the results:

```
do a' <- rpar a  
    b' <- rseq b  
    rseq a'  
    return (a', b')
```

- Alternatively:

```
do a' <- rpar a  
    b' <- rseq b  
    return (a', b')
```

- What is the difference between the two?

## Back to Sudoku

Let's divide the work in two, so we can solve each half in parallel:

```
main :: IO ()
main = do
  [f] <- getArgs
  file <- readFile f
  let puzzles =
        lines file
    (as,bs) =
        splitAt (length puzzles `div` 2) puzzles
  solutions = runEval $ do
    as' <- rpar (map solve as)
    bs' <- rpar (map solve bs)
    return (as' ++ bs')
  print (length (filter isJust solutions))
```

# Back to Sudoku

Let's divide the work in two, so we can solve each half in parallel:

```
main :: IO ()
main = do
  [f] <- getArgs
  file <- readFile f
  let puzzles =
        lines file
        (as,bs) =
          splitAt (length puzzles `div` 2) puzzles
  solutions = runEval $ do
    as' <- rpar (map solve as)
    bs' <- rpar (map solve bs)
    return (as' ++ bs')
  print (length (filter isJust solutions))
```

Does not work!

rpar evaluates its argument to Weak Head Normal Form (WHNF)



# Refresher: Lazy Evaluation

- Evaluating an expression creates a *thunk*


```
map f [1,2,3]
```

# Refresher: Lazy Evaluation

- Evaluating an expression creates a *thunk*

`map f [1,2,3]`

- Weak head normal form (WHNF) evaluates to first constructor




# Refresher: Lazy Evaluation


- Evaluating an expression creates a *thunk*

`map f [1,2,3]`

- Weak head normal form (WHNF) evaluates to first constructor



- Spine evaluated




# Refresher: Lazy Evaluation


- Evaluating an expression creates a *thunk*

`map f [1,2,3]`

- Weak head normal form (WHNF) evaluates to first constructor



- Spine evaluated




# Refresher: Lazy Evaluation


- Evaluating an expression creates a *thunk*

`map f [1,2,3]`

- Weak head normal form (WHNF) evaluates to first constructor



- Spine evaluated



- Stuff like this is why Haskell is not an ideal parallel language.

## We need to go deeper

```
import Control.DeepSeq
```

```
class NFData a where  
  rnf :: a -> ()
```

```
deepseq a b = rnf a 'seq' b
```

```
force :: NFData a => a -> a  
force a = deepseq a a
```

- force fully evaluates a nested data structure and returns it
  - ▶ For example a list: the list is fully evaluated, including the elements. Example from library:

```
instance NFData a => NFData [a] where  
  rnf [] = []  
  rnf (x:xs) = x 'deepseq' xs 'deepseq' ()
```

## Apply force

```
main :: IO ()
main = do
  [f] <- getArgs
  file <- readFile f
  let puzzles =
        lines file
    (as,bs) =
        splitAt (length puzzles `div` 2) puzzles
  solutions = runEval $ do
    as' <- rpar (force (map solve as))
    bs' <- rpar (force (map solve bs))
    return (as' ++ bs')
  print (length (filter isJust solutions))
```

## Apply force

```
main :: IO ()
main = do
  [f] <- getArgs
  file <- readFile f
  let puzzles =
        lines file
    (as,bs) =
        splitAt (length puzzles `div` 2) puzzles
  solutions = runEval $ do
    as' <- rpar (force (map solve as))
    bs' <- rpar (force (map solve bs))
    return (as' ++ bs')
  print (length (filter isJust solutions))
```



# Compile and run the program

```
$ ghc -O2 sudoku2.hs -rtsopts -threaded
[1 of 2] Compiling Sudoku          ( Sudoku.hs, Sudoku.o )
[2 of 2] Compiling Main              ( sudoku2.hs, sudoku2.o )
Linking sudoku2 ...
$ ./sudoku2 sudoku17.1000.txt +RTS -s -N2
  2,383,492,264 bytes allocated in the heap
  49,765,200 bytes copied during GC
  2,489,872 bytes maximum residency (8 sample(s))
  259,184 bytes maximum slop
    9 MB total memory in use (0 MB lost due to fragmentation)

Tot time (elapsed)
Gen  0      2979 colls,  2978 par    0.12s   0.09s
Gen  1         8 colls,    8 par    0.02s   0.01s

SPARKS: 2 (1 converted, 0 overflowed, 0 dud, 0 GC'd, 1 fizzled)

INIT    time    0.00s ( 0.00s elapsed)
MUT     time    2.36s ( 1.43s elapsed)
GC       time    0.14s ( 0.10s elapsed)
EXIT    time    0.00s ( 0.00s elapsed)
Total   time    2.50s ( 1.53s elapsed)
```

# Calculating speedup

- Calculating speedup with 2 processors:
  - ▶ Elapsed time (1 proc) / Elapsed Time (2 procs)
  - ▶ **NOT** CPU time (2 procs) / Elapsed (2 procs)
  - ▶ Compare against sequential program, not parallel program running on 1 CPU (why?)
- Speedup for sudoku2:  $2.11/1.53 = 1.38$

# Why not 2?

There are two main reasons for lack of parallel speedup

1. Less than 100% *utilisation* (some processors idle for part of the time)
2. Extra *overhead* in the parallel version

Each of these has many possible causes.

# A menu of ways to screw up

## 1. Less than 100% utilisation

- ▶ Parallelism was not created, or was discarded
- ▶ Algorithm not fully parallelised residual sequential computation
- ▶ Uneven workloads
- ▶ Poor scheduling
- ▶ Communication latency

## 2. Extra overhead in the parallel version

- ▶ Overheads from rpar, work-stealing, deep, ...
- ▶ Lack of locality, cache effects...
- ▶ Larger memory requirements leads to GC overhead
- ▶ GC synchronisation
- ▶ Duplicating work (but this might increase utilisation!)

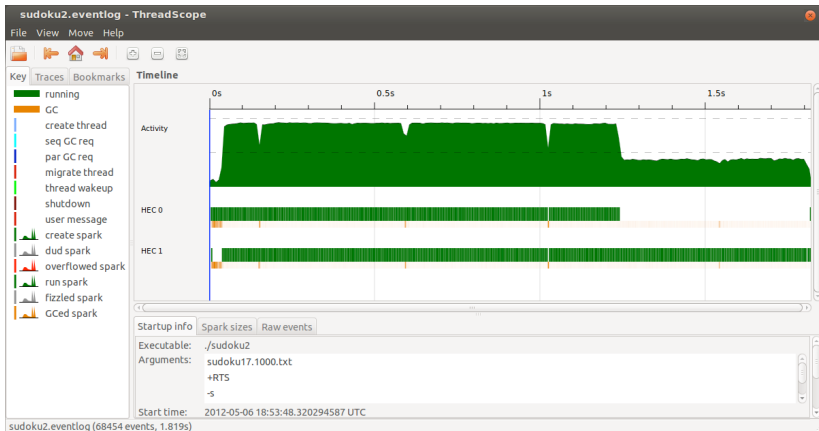
## So we need tools

- We need tools to tell us why the program isn't performing as well as it could be
- For Parallel Haskell we have ThreadScope

## So we need tools

- We need tools to tell us why the program isn't performing as well as it could be
- For Parallel Haskell we have ThreadScope

```
$ ghc -O2 sudoku2.hs -rtsopts -threaded -eventlog  
$ ./sudoku2 sudoku17.1000.txt +RTS -N2 -ls  
$ threadscope sudoku2.eventlog
```



# Uneven workloads

- One of the tasks took longer than the other, leading to less than 100% utilisation

```
let (as, bs) =  
    splitAt (length puzzles 'div' 2)  
           puzzles
```

- One of these lists contains more work than the other, even though they have the same length



# Partitioning

```
let (as, bs) =  
    splitAt (length puzzles `div` 2)  
           puzzles
```

- Dividing up the work along fixed pre-defined boundaries like this example, is called *static partitioning*
- Static partitioning is simple, but can lead to under-utilisation if the tasks can vary in size
- Static partitioning does not adapt to varying availability of processors
- Our solution in this example can use only 2 processors

# Dynamic Partitioning

- *Dynamic partitioning* involves
  - ▶ dividing the work into smaller units
  - ▶ assigning work units to processors dynamically at runtime using a scheduler
- Good for irregular problems and varying number of processors
- GHC's runtime system provides spark pools to track the work units, and a work-stealing scheduler to assign them to processors
- So all we need to do is use smaller tasks and more rpars, and we get dynamic partitioning

## Back to Sudoku

```
runEval $ do  
  as' <- rpar (force (map solve as))  
  bs' <- rpar (force (map solve bs))
```

We want to push `rpar` down into the `map`, so that each call to `solve` will be a separate spark

# A parallel map

```
parMap :: (a -> b) -> [a] -> Eval [b]
parMap f [] = return []
parMap f (a:as) = do
  b <- rpar (f a)
  bs <- parMap f as
  return (b:bs)
```

(Already provided in `Control.Parallel.Strategies`)

## Putting it together

```
main :: IO ()
main = do
  [f] <- getArgs
  file <- readFile f

  let puzzles    = lines file
      solutions = runEval (parMap solve puzzles)

  print (length (filter isJust solutions))
```

## Result with 2 cores

```
$ ./sudoku3 sudoku17.1000.txt +RTS -N2 -s
2,386,461,968 bytes allocated in the heap
68,963,504 bytes copied during GC
2,717,872 bytes maximum residency (14 sample(s))
179,824 bytes maximum slop
9 MB total memory in use (0 MB lost due to fragmentation)
```

			Tot time (elapsed)	
Gen 0	2511 colls,	2510 par	0.14s	0.10s
Gen 1	14 colls,	14 par	0.05s	0.03s

Parallel GC work balance: 1.69 (8603353 / 5095344, ideal 2)

SPARKS: 1000 (1000 converted, 0 overflowed, 0 dud, 0 GC'd, 0 fizzled)

INIT	time	0.00s	( 0.00s elapsed)
MUT	time	2.46s	( 1.24s elapsed)
GC	time	0.19s	( 0.12s elapsed)
EXIT	time	0.00s	( 0.00s elapsed)
Total	time	2.65s	( 1.37s Elapsed)

## Result with 2 cores

```
$ ./sudoku3 sudoku17.1000.txt +RTS -N2 -s
2,386,461,968 bytes allocated in the heap
68,963,504 bytes copied during GC
2,717,872 bytes maximum residency (14 sample(s))
179,824 bytes maximum slop
9 MB total memory in use (0 MB lost due to fragmentation)
```

			Tot time (elapsed)	
Gen 0	2511 colls,	2510 par	0.14s	0.10s
Gen 1	14 colls,	14 par	0.05s	0.03s

Parallel GC work balance: 1.69 (8603353 / 5095344, ideal 2)

SPARKS: 1000 (1000 converted, 0 overflowed, 0 GC'd, 0 fizzled)

INIT	time	0.00s	( 0.00s elapsed)
MUT	time	2.46s	( 1.24s elapsed)
GC	time	0.19s	( 0.12s elapsed)
EXIT	time	0.00s	( 0.00s elapsed)
Total	time	2.65s	( 1.37s Elapsed)

Speedup 1.54

# Evaluation Strategies

- So far we have used Eval/rpar/rseq
  - ▶ These are quite low-level tools
  - ▶ But its important to understand how the underlying mechanisms work
- Now: raise the level of abstraction
- Goal: encapsulate parallel idioms as re-usable components that can be composed together.



## Two High-level Approaches

- Composing Strategy's in the Eval monad, and exploit modularity via lazy evaluation (see the book)
- Use the Par monad to make parallel computations explicit (the rest of the lecture)

# The Par monad

- Abandon modularity via lazy evaluation
- Get a more direct programming model
- Avoid some common pitfalls
- Modularity via higher-order skeletons
- A beautiful implementation

# The Par monad

```
import Control.Monad.Par
```

```
data Par
```

```
instance Monad Par
```

```
runPar :: Par a -> a
```

```
fork :: Par () -> Par ()
```

```
data IVar
```

```
new :: Par (IVar a)
```

```
get :: IVar a -> Par a
```

```
put :: NFData a => IVar a -> a -> Par ()
```

# The Par monad

```
import Control.Monad.Par
```

```
data Par
```

```
instance Monad Par
```

```
runPar :: Par a -> a
```

```
fork :: Par () -> Par ()
```

DANGER!

Must not contain an IVar

```
data IVar
```

```
new :: Par (IVar a)
```

```
get :: IVar a -> Par a
```

```
put :: NFData a => IVar a -> a -> Par ()
```

# Sudoku is addictive

Let's use the Par monad to solve some sudoku puzzles in parallel

```
main :: IO ()
main = do
  [f] <- getArgs
  grids <- fmap lines (readFile f)
  let (as,bs) = splitAt (length grids `div` 2)
                        grids
      solutions = runPar $ do
        ...
        return (as' ++ bs')
  print (length (filter isJust solutions))
```

# Building new abstractions

- Par primitives can be used for building new abstractions.
- For instance, spawning a computation and make a new IVar for the result is often useful (aka futures).

```
spawn :: NFData a => Par a -> Par (IVar a)
spawn p = do r <- new
           fork $ p >>= put r
           return r
```

## Divide and conquer parallelism

```
parfib :: Int -> Int -> Par Int
parfib n
  | n <= 2 = return 1
  | otherwise = do
    x <- spawn $ parfib (n-1)
    y <- spawn $ parfib (n-2)
    x' <- get x
    y' <- get y
    return $ x' + y'
```

## Another abstraction

- We can use `spawn` to build other primitives for regular parallelism, like `parMap`:



## Another abstraction

- We can use `spawn` to build other primitives for regular parallelism, like `parMap`:
- `parMapM :: NFData b`  
     $\Rightarrow (a \rightarrow \text{Par } b) \rightarrow [a] \rightarrow \text{Par } [b]$   
`parMapM f as = do`  
    `ibs <- mapM (spawn . f) as`  
    `mapM get ibs`

## Another abstraction

- We can use `spawn` to build other primitives for regular parallelism, like `parMap`:
- `parMapM :: NFData b`  
     $\Rightarrow (a \rightarrow \text{Par } b) \rightarrow [a] \rightarrow \text{Par } [b]$   
`parMapM f as = do`  
    `ibs <- mapM (spawn . f) as`  
    `mapM get ibs`
- Similar to `parMap`, except that the type is subtly different


# How did we avoid laziness?

- `put` is *hyperstrict*.
- (By default)
- There's also a WHNF version called `put_`
- **Why might we want to use `put_`?**

# Dataflow problems

- `Par` really shines when the problem is easily expressed as a dataflow graph, particularly an irregular or dynamic graph (e.g., shape depends on the program input)
- Identify the nodes and edges of the graph
  - ▶ Each node is created by `fork`
  - ▶ Each edge is an `IVar`
- Many problems that are awkward in e.g. Futhark are straightforward with `Par`

## Par expresses dynamic dataflow



## Thoughts to take away

- Parallelism is not the goal
- Making your program *faster* is the goal (unlike Concurrency, which is a goal in itself)
- If you can make your program fast enough without parallelism, all well and good

- On Wednesday: Work with parallelism in Haskell
- Assignment:
  - ▶ Speeding up a ray tracer by computing pixels in parallel
  - ▶ Speeding up a ray tracer by computing *BVH* in parallel
  - ▶ Maybe other sources of parallelism, if you can find them!